

**Per-decision inference cost: EmRL is flat, classical grows with size**

