

Scalability vs candidate window: widening K recovers delivery at scale (no retraining — K-agnostic network)

